

# Import libraries

```
!pip install pmdarima
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import tensorflow as tf
import os
from sklearn.preprocessing import MinMaxScaler
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.seasonal import seasonal_decompose
from pmdarima.arma import auto_arma
from statsmodels.tsa.arma.model import ARIMA
```

```
import warnings
warnings.filterwarnings('ignore')
```

```
Requirement already satisfied: pmdarima in
/usr/local/lib/python3.10/dist-packages (2.0.4)
Requirement already satisfied: joblib>=0.11 in
/usr/local/lib/python3.10/dist-packages (from pmdarima) (1.4.2)
Requirement already satisfied: Cython!=0.29.18,!=0.29.31,>=0.29 in
/usr/local/lib/python3.10/dist-packages (from pmdarima) (3.0.11)
Requirement already satisfied: numpy>=1.21.2 in
/usr/local/lib/python3.10/dist-packages (from pmdarima) (1.26.4)
Requirement already satisfied: pandas>=0.19 in
/usr/local/lib/python3.10/dist-packages (from pmdarima) (2.2.3)
Requirement already satisfied: scikit-learn>=0.22 in
/usr/local/lib/python3.10/dist-packages (from pmdarima) (1.2.2)
Requirement already satisfied: scipy>=1.3.2 in
/usr/local/lib/python3.10/dist-packages (from pmdarima) (1.13.1)
Requirement already satisfied: statsmodels>=0.13.2 in
/usr/local/lib/python3.10/dist-packages (from pmdarima) (0.14.4)
Requirement already satisfied: urllib3 in
/usr/local/lib/python3.10/dist-packages (from pmdarima) (2.3.0)
Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in
/usr/local/lib/python3.10/dist-packages (from pmdarima) (75.1.0)
Requirement already satisfied: packaging>=17.1 in
/usr/local/lib/python3.10/dist-packages (from pmdarima) (24.2)
Requirement already satisfied: mkl_fft in
/usr/local/lib/python3.10/dist-packages (from numpy>=1.21.2->pmdarima)
(1.3.8)
Requirement already satisfied: mkl_random in
/usr/local/lib/python3.10/dist-packages (from numpy>=1.21.2->pmdarima)
(1.2.4)
Requirement already satisfied: mkl_umath in
/usr/local/lib/python3.10/dist-packages (from numpy>=1.21.2->pmdarima)
```

```

(0.1.1)
Requirement already satisfied: mkl in /usr/local/lib/python3.10/dist-
packages (from numpy>=1.21.2->pmdarima) (2025.0.1)
Requirement already satisfied: tbb4py in
/usr/local/lib/python3.10/dist-packages (from numpy>=1.21.2->pmdarima)
(2022.0.0)
Requirement already satisfied: mkl-service in
/usr/local/lib/python3.10/dist-packages (from numpy>=1.21.2->pmdarima)
(2.4.1)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.10/dist-packages (from pandas>=0.19->pmdarima)
(2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.10/dist-packages (from pandas>=0.19->pmdarima)
(2025.1)
Requirement already satisfied: tzdata>=2022.7 in
/usr/local/lib/python3.10/dist-packages (from pandas>=0.19->pmdarima)
(2025.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.22-
>pmdarima) (3.5.0)
Requirement already satisfied: patsy>=0.5.6 in
/usr/local/lib/python3.10/dist-packages (from statsmodels>=0.13.2-
>pmdarima) (1.0.1)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2-
>pandas>=0.19->pmdarima) (1.17.0)
Requirement already satisfied: intel-openmp>=2024 in
/usr/local/lib/python3.10/dist-packages (from mkl->numpy>=1.21.2-
>pmdarima) (2024.2.0)
Requirement already satisfied: tbb==2022.* in
/usr/local/lib/python3.10/dist-packages (from mkl->numpy>=1.21.2-
>pmdarima) (2022.0.0)
Requirement already satisfied: tcmlib==1.* in
/usr/local/lib/python3.10/dist-packages (from tbb==2022.*->mkl-
>numpy>=1.21.2->pmdarima) (1.2.0)
Requirement already satisfied: intel-cmplr-lib-rt in
/usr/local/lib/python3.10/dist-packages (from mkl_umath-
>numpy>=1.21.2->pmdarima) (2024.2.0)
Requirement already satisfied: intel-cmplr-lib-ur==2024.2.0 in
/usr/local/lib/python3.10/dist-packages (from intel-openmp>=2024->mkl-
>numpy>=1.21.2->pmdarima) (2024.2.0)

```

## Load the Dataset

```

df =
pd.read_csv('/kaggle/input/stock-market-dataset/stocks/NVDA.csv', parse

```

```
_dates=['Date'])
df.head()
```

	Date	Open	High	Low	Close	Adj Close
Volume						
0	1999-01-22	1.750000	1.953125	1.552083	1.640625	1.509998
67867200.0						
1	1999-01-25	1.770833	1.833333	1.640625	1.812500	1.668188
12762000.0						
2	1999-01-26	1.833333	1.869792	1.645833	1.671875	1.538759
8580000.0						
3	1999-01-27	1.677083	1.718750	1.583333	1.666667	1.533965
6109200.0						
4	1999-01-28	1.666667	1.677083	1.651042	1.661458	1.529172
5688000.0						

## Basic Analysis

```
print('Basic information about the data:',df.info())
print('\n\nNull values:\n', df.isna().sum())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5334 entries, 0 to 5333
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date        5334 non-null   datetime64[ns]
1   Open        5333 non-null   float64
2   High        5333 non-null   float64
3   Low         5333 non-null   float64
4   Close       5333 non-null   float64
5   Adj Close   5333 non-null   float64
6   Volume      5333 non-null   float64
dtypes: datetime64[ns](1), float64(6)
memory usage: 291.8 KB
Basic information about the data: None
```

```
Null values:
Date      0
Open      1
High      1
Low       1
Close     1
Adj Close 1
Volume    1
dtype: int64
```

## Drop null values

```
df = df.dropna()  
df.isna().sum()
```

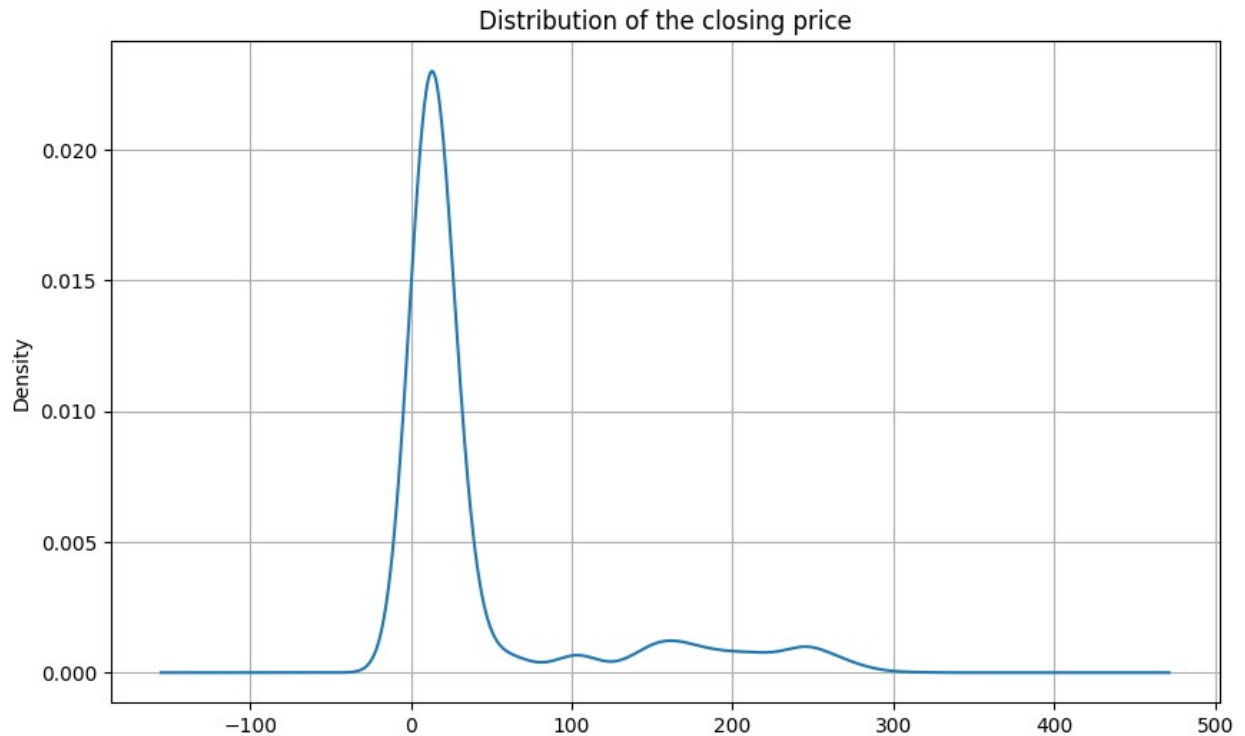
```
Date          0  
Open          0  
High          0  
Low           0  
Close         0  
Adj Close     0  
Volume        0  
dtype: int64
```

## Sessional Trend of the closing price

```
# plot close price  
plt.figure(figsize=(10,6))  
plt.grid(True)  
plt.xlabel('Date')  
plt.ylabel('Close Prices')  
plt.plot(df['Close'])  
plt.title('Closing Price')  
plt.show()
```



```
plt.figure(figsize=(10,6))
df_close = df['Close']
df_close.plot(kind='kde', title = 'Distribution of the closing price')
plt.grid(True)
plt.show()
```




```
# stationary test
def test_stationarity(timeseries):
    rolmean = timeseries.rolling(12).mean()
    rolstd = timeseries.rolling(12).std()
    plt.figure(figsize=(12,6))
    plt.plot(timeseries,color='blue',label='Original')
    plt.plot(rolmean,color='red',label='rolling mean')
    plt.plot(rolstd,color='black',label='rolling std')
    plt.legend(loc='best')
    plt.title('Rolling Mean and Standard Deviation')
    plt.show(block=False)
    print('Dickey Fuller Test')
    adft = adfuller(timeseries,autolag='AIC')

    output = pd.Series(adft[0:4], index=['Test Statistics', 'p-
value','No. of lags used', 'Number of observations used'])

    for key, values in adft[4].items():
        output['critical value (%s)'%key] = values

    print(output)

test_stationarity(df_close)
```




```
Dicky Fuller Test
Test Statistics          0.214695
p-value                 0.973074
No. of lags used        32.000000
Number of observations used 5300.000000
critical value (1%)     -3.431584
critical value (5%)     -2.862085
critical value (10%)    -2.567060
dtype: float64
```

## Eliminate Trend and Seasonality

```
result = seasonal_decompose(df_close,model='multiplicative',period=30)
fig = plt.figure()
fig = result.plot()
fig.set_size_inches(16,9)
plt.show()
```


<Figure size 640x480 with 0 Axes>



```
#Eliminate trend
from pylab import rcParams

rcParams['figure.figsize'] = 10, 6
df_log = np.log(df_close)
moving_avg = df_log.rolling(12).mean()
std_dev = df_log.rolling(12).std()
plt.legend(loc='best')
plt.title('Moving Average')
plt.plot(std_dev, color="black", label = "Standard Deviation")
plt.plot(moving_avg, color="red", label = "Mean")
plt.legend()
plt.grid(True)
plt.show()
```






## Split data into training and testing part

*#split data into train and training set*

```
train_data, test_data = df_log[3:int(len(df_log)*0.9)],  
df_log[int(len(df_log)*0.9):]  
plt.figure(figsize=(10,6))  
plt.grid(True)  
plt.xlabel('Dates')  
plt.ylabel('Closing Prices')  
plt.plot(df_log, 'green', label='Train data')  
plt.plot(test_data, 'blue', label='Test data')  
plt.legend()  
plt.show()
```



## Auto ARIMA Find the Best Parameters

```
model_autoARIMA = auto_arima(train_data, start_p=0, start_q=0,
                              test='adf',          # use adftest to find optimal
                              'd'
                              max_p=3, max_q=3, # maximum p and q
                              m=1,             # frequency of series
                              d=None,          # let model determine 'd'
                              seasonal=False,  # No Seasonality
                              start_P=0,
                              D=0,
                              trace=True,
                              error_action='ignore',
                              suppress_warnings=True,
                              stepwise=True)
print(model_autoARIMA.summary())
model_autoARIMA.plot_diagnostics(figsize=(15,8))
plt.show()
```

Performing stepwise search to minimize aic

```
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-17432.104, Time=0.86 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=-17431.545, Time=0.47 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=-17431.589, Time=0.39 sec
ARIMA(0,1,0)(0,0,0)[0]          : AIC=-17430.730, Time=0.32 sec
```

ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=-17428.141, Time=1.79 sec

Best model: ARIMA(0,1,0)(0,0,0)[0] intercept

Total fit time: 3.869 seconds

#### SARIMAX Results

Dep. Variable: y No. Observations: 4796

Model: SARIMAX(0, 1, 0) Log Likelihood 8718.052

Date: Wed, 09 Apr 2025 AIC -17432.104

Time: 06:59:31 BIC -17419.153

Sample: 0 HQIC -17427.555

- 4796

Covariance Type: opg

	coef	std err	z	P> z	[0.025
--	------	---------	---	------	--------

0.975]

intercept	0.0010	0.001	1.836	0.066	-7.04e-05
-----------	--------	-------	-------	-------	-----------

sigma2	0.0015	1.15e-05	134.629	0.000	0.002
--------	--------	----------	---------	-------	-------

Ljung-Box (L1) (Q): 1.44 Jarque-Bera (JB): 34511.77

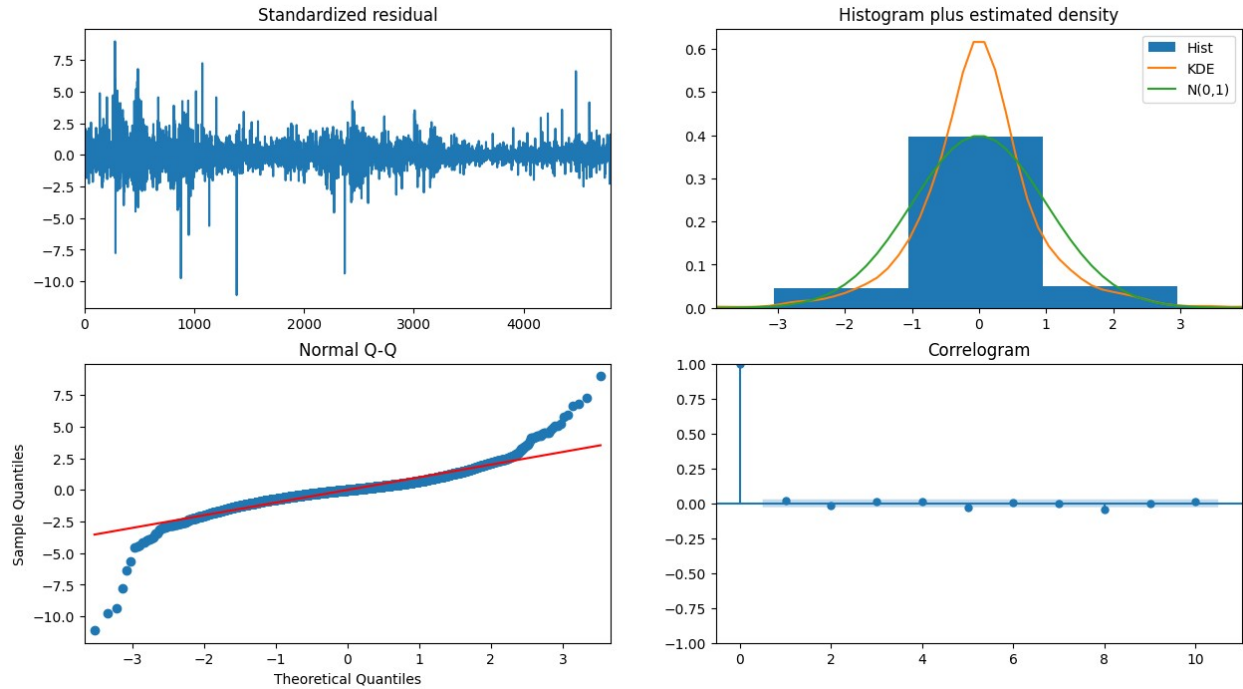
Prob(Q): 0.23 Prob(JB): 0.00

Heteroskedasticity (H): 0.17 Skew: -0.14

Prob(H) (two-sided): 0.00 Kurtosis: 16.14

#### Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).



## Modeling

```
#Modeling
# Build Model
model = ARIMA(train_data, order=(1,1,2))
fitted = model.fit()
print(fitted.summary())
```

### SARIMAX Results

```
=====
=====
Dep. Variable:                  Close    No. Observations:
4796
Model:                        ARIMA(1, 1, 2)    Log Likelihood
8717.191
Date:                        Wed, 09 Apr 2025    AIC
17426.382
Time:                        07:03:35    BIC
17400.480
Sample:                        0    HQIC
17417.285
- 4796
Covariance Type:                opg
=====
```

```

=====
              coef      std err          z      P>|z|      [0.025
0.975]
-----
-----
ar.L1          -0.7227      1.503      -0.481      0.631      -3.668
2.223
ma.L1           0.7412      1.501       0.494      0.621      -2.200
3.682
ma.L2           0.0112       0.032       0.350      0.726      -0.052
0.074
sigma2          0.0015     1.15e-05     133.814      0.000       0.002
0.002
=====
=====
Ljung-Box (L1) (Q):                0.00   Jarque-Bera (JB):
34690.33
Prob(Q):                          0.95   Prob(JB):
0.00
Heteroskedasticity (H):            0.17   Skew:
-0.14
Prob(H) (two-sided):              0.00   Kurtosis:
16.17
=====
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients
(complex-step).

```

## Modeling and Forecasting

```

# Forecast
forecast_result = fitted.get_forecast(steps=321)
fc = forecast_result.predicted_mean
conf = forecast_result.conf_int(alpha=0.05)

# Make as pandas series
fc_series = pd.Series(fc, index=test_data.index)
lower_series = pd.Series(conf.iloc[:, 0], index=test_data.index)
upper_series = pd.Series(conf.iloc[:, 1], index=test_data.index)


# Plot
plt.figure(figsize=(10,5), dpi=100)
plt.plot(train_data, label='Training Data')
plt.plot(test_data, color='blue', label='Actual Stock Price')
plt.plot(fc_series, color='orange', label='Predicted Stock Price')
plt.fill_between(lower_series.index, lower_series, upper_series,

```

```

        color='k', alpha=0.1)
plt.title('ARCH CAPITAL GROUP Stock Price Prediction')
plt.xlabel('Time')
plt.ylabel('ARCH CAPITAL GROUP Stock Price')
plt.legend(loc='upper left', fontsize=8)
plt.show()

```



## Evaluate Model Performance

```

# Align test data to forecast length
aligned_test = test_data[:len(fc)]

# Then compute metrics
mse = mean_squared_error(aligned_test, fc)
mae = mean_absolute_error(aligned_test, fc)
rmse = math.sqrt(mse)
mape = np.mean(np.abs(fc - aligned_test) / np.abs(aligned_test)) * 100

# Print results
print(f"MSE: {mse:.4f}")
print(f"MAE: {mae:.4f}")
print(f"RMSE: {rmse:.4f}")
print(f"MAPE: {mape:.2f}%")

MSE: 0.0882
MAE: 0.2178

```

RMSE: 0.2970  
MAPE: 4.18%

**'with a MAPE of around 4.18% the model is 95.82 % accurate in predicting the next observations**